

---

# Neural Architecture Search: An Overview

---

**Federico Minutoli**

DIBRIS

University of Genoa

fede97.minutoli@gmail.com

**Massimiliano Ciranni**

DIBRIS

University of Genoa

massimiliano.ciranni@outlook.com

## Abstract

Recent advances in deep learning research have given birth to really powerful architectures, which have been used with great success in a wide range of fields. Despite designing neural architectures has become easier, a lot of human expertise is still required to find optimal solutions and often comes as the result of extended trial and error cycles. This has led to a rapid increase in the complexity of architectures (e.g., ResNet with its many skip connections) that has left a large segment of the machine learning community wondering "*How can someone come out with such an architecture?*" (J. Langford referring to *DenseNet*) [1]. This is a hot-topic in machine learning right-now, which is being addressed by neural architecture search (NAS), a novel branch of automated machine learning that aims at automating architecture engineering. The core idea behind a NAS system is to feed it with a dataset and a task (classification, regression, etc.) as input, expecting it to find an architecture by navigating through a search space of all possible architectures trying to maximize a reward-based search strategy. In this paper we will provide an overview of the many questions that NAS tries to answer, also addressing some of the main limitations that plague it at the moment (i.e., GPU time, combinatorial explosion, domain shift, etc.). Then we will provide a brief outline of the main approaches followed by a more in depth analysis on two of them: differentiable search spaces (DNAS) and (efficient) reinforcement learning (ENAS). In the end, we will present a naive toy implementation of a controller-based NAS system tailored at sampling multi-layer perceptrons (MLPs) on a handful of datasets, against Scikit-learn's pre-made MLP classifier.

## 1 Introduction

The success that deep learning gained in the last few years is mostly due to its capability of automating the feature engineering process, allowing to extract hierarchical features in an end-to-end fashion, rather than manually design features from data. New opportunities made available by deep learning methods imply that there is a high demand of deep learning expertise and architecture engineering, as we see models becoming more and more complex, while still manually designed. Deep learning engineers are expected to have an intuitive understanding of what architecture might work best for what situation, but this is rarely the case. DNNs are complex systems, and the number of architectural variants one could come up with are literally infinite; consider for example a famous image recognition model, such as ResNet [2]: it comes with many flavours in terms of the number of hidden layers (50, 101, 152) and configurations for its skip-connections. It is obvious that there are many ways in which a model like this could be tweaked just by adding/removing convolutional blocks or skip-connections and, in general, when we look at the best performing models, it is almost never clear why certain design choices have been made in that specific manner. A logical step to avoid the process of manual designing neural networks is to automate it by building a system that given a dataset provides a neural network architecture that can be implemented, trained, evaluated and eventually serve its purpose; this is something that, in comparison, has been already addressed and

solved for classical machine learning algorithms (under the name of hyper-parameter optimization [3]), by means of grid search, random search, *Bayesian optimization*, meta-learning, to name a few, but when considering deep learning architectures, the problem becomes much harder to deal with. The process of automating architecture engineering took the name of *Neural Architecture Search (NAS)*, a branch of automated machine learning (AutoML [4]) which was first introduced by Zoph et al. (2016) [5] and quickly turned into an extremely popular area of research, with a seemingly exponential increase in the number of papers written on the subject, as Figure 1 shows. On top of that all the main manufacturers seem pretty active at the moment, with the aim of uncovering the full potential of such systems, like Google for large-scale object classification [6], Tesla for automatic configuration assessment in automotive [7] and Microsoft with its NNI’s toolkit [8] which offers implementations for most of the current state-of-the-art approaches in NAS.

In the next sections you will find a comprehensive overview of the main components that compose a NAS system and its main limitations (Section 2), a detailed description of two algorithms, DARTS and ENAS, which have received a lot of attention recently (Section 3), followed by our proposed naive implementation of a controller-based NAS system for multi-layer perceptrons (Section 4) and, finally, a few noticeable results that we could gather out of it (Section 5.1) and conclusions (Sections 6).




Figure 1: NAS papers per year based on the literature list on [automl.org](http://automl.org).<sup>2</sup>

## 2 Problem setting analysis

A NAS system can be defined in terms of three main components (see Figure 2)

- *Search space*, shapes the type of valid networks that can be designed.
- *Search strategy*, defines the approach to navigate through the search space.
- *Performance estimation strategy*, looks for ways to evaluate the performance of a plausible network solely from its design without constructing or training it.




Figure 2: Scheme of a generic NAS system. [9]

### 2.1 Search space

The search space defines which neural architectures a NAS approach might discover in principle, that is, the domain  $\mathcal{D}$  of valid architectures it could navigate towards (see an example in Figure 3). A relatively simple search space is the space of *chain-structured neural networks*, a sequence of  $N$  layers, where the  $i$ -th layer  $L_i$  receives its input from the  $(i-1)$ -th layer and sends it to the  $(i+1)$ -th through a function  $g(L_{i-1}^{out})$ . The search space is then parametrized by: (i) the (maximum) number of layers  $N$  (possibly unbounded); (ii) the type of operation every layer executes, e.g.,

<sup>2</sup>The literature list on NAS papers at [www.automl.org](http://www.automl.org) is manually curated and lists all NAS paper that we are aware of starting from early 2017 up to summer 2020.

<sup>3</sup>Implied in Zoph et al. (2016) first NAS system [5]




Figure 3: The search space for CNNs design <sup>3</sup>

pooling, convolution, or more advanced operations like depth-wise separable convolution; and (iii) hyper-parameters associated with the operation, e.g., number of filters, kernel size and strides for a convolutional layer or even number of units for fully-connected networks. Note that the parameters from (iii) are conditioned on (ii), hence the parametrization of the search space is not fixed-length, but a conditional space. Recent work on NAS (Real et al., 2019) incorporates modern design elements known from hand-crafted architectures, such as skip connections, which allow to build complex, *multi-branch networks* (i.e., residual networks, dense networks, etc.). In this case the input of layer  $i$  can be instead formally described as a function  $g(L_{i-1}^{out}, \dots, L_0^{out})$  combining previous layers outputs, in turn resulting in a function with significantly more degrees of freedom. The motif of these approaches is to predict an architecture without previous knowledge of what a good architecture should even look like to begin with, which opens up room to novel architectures and to a better interpretability of the current human-made ones; hence, they’ve been given the name of *macro-architecture* search NAS.

The biggest drawback of the approach mentioned above was the time it took to navigate through the search space before coming up with a definite solution. Zoph et al. (2016) had to use 800 GPUs for 28 days (with a cost of approx. 130k US\$) to navigate through the entire search space before coming up with the best architecture. There was clearly a need for a way to design controllers that could navigate the search space more intelligently and, motivated by hand-crafted architectures consisting of repeated motifs, Zoph et al. (2018) proposed to search for such motifs, dubbed *cells*, rather than for whole architecture. A typical cell is modeled as a *Direct Acyclic Graph* (DAG) consisting of  $n$  nodes. Each node is a latent representation (like 2D kernels in CNNs) while directed edges connecting each node map to operations, such as pooling, or batch normalization.

Zoph et al. (2018) would rather optimize two different kind of cells: a normal cell that preserves the dimensionality of the input and a reduction cell which reduces the spatial dimension. The final architecture is then built by stacking these cells in a predefined manner, dictated by human bias. This approach has brought great benefits to NAS systems research, since the search space of cells consists of significantly less layers than whole architectures, with an estimated speed-up of a factor 7, coming from an order of magnitude less parameters. Also, architectures built from cells can more easily be transferred or adapted to other data sets by simply varying the number of cells and filters used within a model. Indeed, Zoph et al. (2018) transferred cells optimized on CIFAR-10 to ImageNet and achieved state-of-the-art performance, whereas domain shift, due to an intrinsic bias towards the current dataset, has always been a worry-some problem for macro-architectures NAS systems. A few questions rose, though, regarding how while using a cell-based, or *micro-architecture*, search space the macro-architecture would be chosen: how many cells shall be used and how should they be connected in the model? Ideally, both the macro-architecture and the micro-architecture should be optimized jointly instead of solely optimizing the micro-architecture. Until recently this proposition has been omitted, as given the undeniable advantages micro-architecture has shown over macro-architecture, a lot of the community has just shifted towards designing such systems, instead of trying to understand why they outperform macro-architecture ones on many different tasks. Indeed starting from [10], more cost-efficient macro-architecture NAS systems have been designed with the intention of bridging the gap between these two opposite flavours.

## 2.2 Search strategy

The search strategy details how to explore the search space (which often happens to be exponentially large or even unbounded). A good search strategy should be able to balance the trade-off between finding well-performing architectures in short time and premature convergence to a region of sub-

optimal architectures, whose importance is so keen, that most of the work gone into neural architecture search has been innovations for this part of the problem: finding out which optimization methods work best, and how they can be changed or tweaked to make the search process churn out better results faster and with consistent stability. There have been several approaches attempted to navigate the space of neural architectures, including random search, gradient-based methods, Bayesian optimization, reinforcement learning (RL) and evolutionary methods.

### 2.2.1 Reinforcement learning

Reinforcement learning is the science of optimal decision making. When an infant plays, waves its arms, it has no explicit teacher, but it does have a direct sensor motor connection to its environment. Exercising this connection produces a wealth of information about cause and effect, about consequences of actions, and about what to do in order to achieve goals. This is the key idea behind RL: given an environment which represents the outside world to the agent and an agent that takes actions, receives observations from the environment that consists of a reward for his action and information of his new state. That reward informs the agent of how good or bad was the taken action, and the observation tells him what is his next state in the environment. The agent tries to figure out the best actions to take or the optimal way to behave in the environment in order to carry out his task in the best possible way, essentially functioning as an active learning paradigm. Defined as such, in RL there is no concept of supervision. Feedback from the environment might be delayed over several time steps, it's not necessarily instantaneous. This means that the data in RL cannot be consider as i.i.d since the agent might spend some time exploring and wandering in the environment and not see other parts which might be interesting to learn the optimal behavior. So time really matters, the agent must explore pretty much every part of the environment to be able to take the right actions. All of this is formally explained in terms of a policy  $\pi$ , regulated by a parameter  $\theta$ , according to which the agent tries to find the optimal trajectory  $\tau$  across the environment he is in by taking a sequence of actions  $[a_1, \dots, a_T]$  from its corresponding current state  $[s_1, \dots, s_T]$ . Hence, the objective of RL is to maximize the policy gradient w.r.t.  $\theta$ , according to the following gradient descent:

$$\theta_{t+1} = \theta_t + \nabla_{\theta} J(\theta)$$

where  $J(\theta)$  can be expressed in terms of an expectation on future rewards, taking into account the sum of probabilities of the trajectory  $\tau$  conditioned by  $\theta$  and the current rewards at time  $t$ :

$$J(\theta) = \mathbb{E}[\sum_{t=1}^T R(a_t|s_t; \pi_{\theta})] = \sum_{\tau} p(\tau; \theta) R(\tau)$$

which turns the objective function into finding the policy (parameter)  $\theta$  that induces a trajectory  $\tau = [(s_1, a_1), \dots, (s_T, a_T)]$  that maximizes the expected rewards:

$$\theta^* = \arg \max_{\theta} J(\theta) = \arg \max_{\theta} \sum_{\tau} p(\tau; \theta) R(\tau)$$

Being such a powerful tool, not to mention the greatest success in AI history, that is DeepMind's AlphaGo ruthlessly managing to defeat Lee Sedol in 2016, the first NAS paper from Zoph et al. (2016) [5] implied the use a recurrent network to generate the model descriptions of neural networks and train this RNN with reinforcement learning to maximize the expected accuracy of the generated architectures on a validation set which rivaled human made state-of-the-art models on the CIFAR-10 dataset' test accuracy. Many RL-based NAS systems have been proposed since then, but the general idea to frame a NAS as reinforcement learning problem is to associate the generation of a neural architecture to be the agent's action, with the action space identical to the search space. The agent's reward is based on an estimate of the performance of the trained architecture on unseen data and the state can be seen as a summary of the actions sampled so far. Different RL approaches differ in how they represent the agent's policy and how they optimize it: Zoph and Le (2017) use a recurrent neural network (RNN) policy to sequentially sample a string that in turn encodes the neural architecture trained on the popular REINFORCE policy gradient algorithm (Williams, 1992). Often, the models built with the sole objective of a high validation accuracy end up being high in complexity—meaning a greater number of parameters, more memory required, and higher inference times. An important

contribution with to combat these issues was made by Hsu et al. (2018), where they proposed MONAS, which introduces a multi-task objective that not only tries to maximize validation accuracy, but also to minimize power consumption.

### 2.2.2 Progressive NAS

Progressive NAS – still based on RL – is designed to address the task of finding the best network cell for a given task (micro-search). The first iteration of a progressive NAS generates all possible cells made of just one block and creates proxy CNNs formed by a single cell stacked for a predetermined number of times  $n$ . Validation accuracies from proxy CNNs and their respective cell structure are used to train a surrogate function implemented through a LSTM-based network, that will learn to predict the performance of a candidate cell structure without the need of building and training it. As iterations follow, blocks are incrementally added to the original 1-block cells in a greedy manner following a simple heuristic space. The performance predictor has the role of pruning out all those newly found structures that are not promising enough, keeping just the top  $K$  cells for proxy CNNs generation and predictor training. When the desired number of per-cell blocks,  $B$ , is reached, the best performing cell is stacked  $n$  times to form the final full CNN.

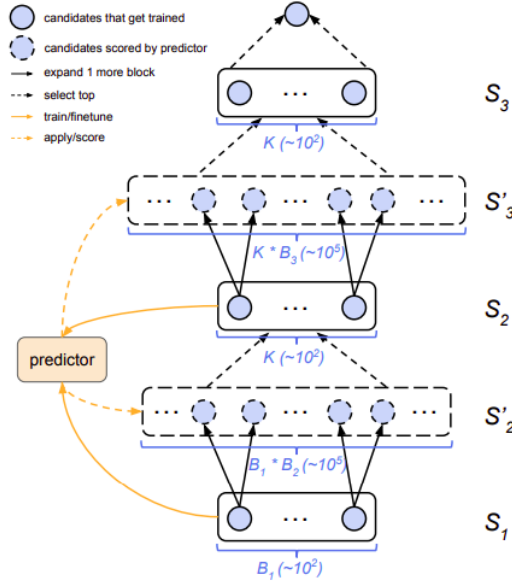


Figure 4: Scheme of the main PNAS steps. [9]

### 2.2.3 Differentiable NAS

Differentiable NAS is a small family of NAS systems which approach the problem of architecture search from a different angle. Instead of searching over a discrete set of candidate architectures, DNAS models leverage a continuous relaxation of the search space, so that the architecture can be optimized with respect to its validation set performance by gradient descent. The data efficiency of gradient-based optimization, as opposed to inefficient black-box search, allows differentiable methods to achieve competitive performance with state-of-the-art models using orders of magnitude less computation resources. For this reason, this family of NAS models is sometimes referred to as *super-network NAS*, as these models were implemented to challenge the main state-of-the-art networks performing complex vision tasks. The idea of searching architectures within a continuous domain was already explored in [11], [12], [13] and [14], but it was only with the introduction of DARTS [15] that this approach gained its current popularity.

### 2.2.4 Others

A strong alternative to using RL are aforementioned neuro-evolutionary approaches that use evolutionary algorithms for optimizing the neural architecture. Evolutionary algorithms evolve a population of models, i.e., a set of (plausible) networks; in every evolution step, at least one model from the population is sampled and serves as a parent to generate off-springs by applying mutations to it. Neuro-evolutionary methods differ in how they sample parents, update populations, and generate offsprings by either applying tournament selection (Goldberg and Deb, 1991) or follow a multi-objective Pareto front to sample parents. Floreano et al. (2008) claimed that gradient-based methods outperform evolutionary methods for the optimization of neural network weights, and that evolutionary approaches should only be used to optimize the architecture itself, yet Stanley et al. (2019) highlighted the ability to account for diversity in genetic algorithms, which increases the chance of finding novel architectures and also allows for massively parallel exploration.

Bayesian optimization is one of the most popular methods for hyperparameter optimization, but it has not been applied to NAS by many groups since typical BO toolboxes are based on Gaussian processes and focus on low-dimensional continuous optimization problems. NASBOT from Kandasamy et al. (2018) is still worth mentioning for its efficiency.

While not exactly in the domain of architecture search, an important consideration in designing neural architectures is that of model compression. Several approaches have been adapted over time: quantization, pruning, knowledge distillation, etc. It is worth mentioning Stier et al. (2019)’s work, who proposed a game theoretic approach based on Shapley values to create efficient topologies. MorphNet [16] by Google follows such topologies, and after being applied on Inception V2 trained on ImageNet, it has scored the same test accuracy, while consuming around 10% less FLOPS.

## 2.3 Performance estimation

The main goal of NAS is to find architectures that achieve high predictive performances on a certain task with the ability to generalize when dealing with unseen data. Performance estimation refers exactly to the process of estimating this performance: the simplest option is to perform a standard training and validation of a generated architecture on data, and measure its performances, but this is unfortunately computationally expensive and limits the number of architectures that can be explored. Much recent research therefore has been focusing on developing methods to address this issue by reducing the cost of these performance estimations. A simple approach relies on low fidelity estimates obtained by training for fewer epochs or on a small subset of data as is done by Real et al. (2019). This approach could work if we could be sure that the relative ranking of architectures does not change due to the low fidelity evaluations, but recent research has shown this not to be the case. Another strategy is based on extrapolating from learning curves: architectures that are predicted to perform poorly from the learning curve in the initial few epochs can be terminated early to speed up the search process. Liu et al. (2018), on the other hand, don’t use learning curves but suggest training a surrogate model (mentioned in 2.2.2), referred to as an *accuracy predictor*, used to predict the performance of an architecture based on properties extrapolated from other novel architectures in an adversarial manner. Other approaches rely on *parameter sharing*, either in the form of network morphism, where child networks are warm-started from surrogate parent networks’ weights, or in the form of one-shot architecture search, popularized by Pham et al. (2018). The latter represents a family of NAS methods which attempt to train one super-architecture which subsumes all the other architectures in the search space. In order to do so parameters are shared between all the networks that are created by the search space navigation strategy, which also allows for fast evaluation since the weights of the super-graph can be inherited to evaluate the various different architectures without training them. While one-shot NAS has reduced the computational resources required for neural architecture search, it is unclear how the inheritance of weights from the super-graph affects the performance of the sub-graphs and the kind of biases it introduces into architectures, as [1] mentioned. Recent explorations on the effect of reduced training in NAS [17] have also discovered high rank correlations between various fully and partially trained connected architectures, whereas others [18] have had success optimizing a NAS in limited GPU time following a gene expression approach.

### 3 Algorithmic overview

#### 3.1 DARTS

In 2019 Liu, Simonyan and Yang proposed DARTS [15], a differentiable architecture search method based on the continuous relaxation of typically discrete NAS search spaces. It also leverages game theoretic concepts for the optimization of neural architectures, as it poses the architecture search problem as a bi-level optimization, where one optimization problem is embedded in another.

DARTS attempts to solve three main problems that are typically encountered in NAS:

- Extremely large search spaces (*micro-search approach*).
- Discrete search space causing expensive optimization (*continuous relaxation*).
- Every candidate architecture must be trained from scratch (*parameter sharing*).

##### 3.1.1 Search space

Following the recent advances in neural architecture search, DARTS embraces micro-search, looking for the best layer configuration instead of taking into consideration a complete multi-layer architecture. As already mentioned in section 2.1, these basic building blocks are defined as *cells*.

A DARTS cell is represented as a DAG, in which each node  $x^{(i)}$  is a latent representation (e.g. a 2D kernel inside a CNN) and each directed edge  $(i, j)$  is associated with some operation  $o^{(i,j)}$ , for example a pooling operation, that transforms  $x^{(i)}$ . Each intermediate node is computed from all of its predecessors, as in eq.(1)

$$x^{(j)} = \sum_{i < j} o^{(i,j)}(x^{(i)}) \quad (1)$$

It is important to point out that a *zero* operation is always included into the search space, as it is needed in order to express the lack of connection between two nodes inside a cell.

Let  $\mathcal{O}$  be a set of candidate operations, where each operation  $o$  represents some function  $o(\cdot)$  to be applied to a node  $x^{(i)}$ . We also denote as  $N$  the number of nodes that we want our cell to be composed of. Given two nodes  $x^{(i)}, x^{(j)}$  the directed edge  $(i, j)$  connecting them is parameterized by a vector  $\alpha^{(i,j)} \in \mathbb{R}^{|\mathcal{O}|}$ .

In order to obtain a continuous search space, the categorical choice of a assigning a particular operation to an edge  $(i, j)$  is relaxed to a *softmax* over all possible operations (2).

$$\bar{o}^{(i,j)}(x) = \sum_{o \in \mathcal{O}} \frac{\exp(\alpha_o^{(i,j)})}{\sum_{o' \in \mathcal{O}} \exp(\alpha_{o'}^{(i,j)})} o(x) \quad (2)$$

This has two major consequences: (i), the operation of an edge is the sum of all the operations placed on that edge, weighted by the *softmax* over  $\alpha^{(i,j)}$ ; (ii), the task of architecture search reduces to learning a set of continuous variables  $\alpha$  so that  $\alpha = \{\alpha^{(i,j)} | i \neq j \wedge i, j \in [1, N]\}$ , from which probability distributions of what operations are to be chosen for each edge are learned. Therefore, *finding the optimal cell* is to find the best operation to assign on each edge. The great advantage of DARTS is that each different choice is not investigated independently, but every possible operation is mixed so their weights are actually shared and "trained together". When the search ends, a discrete architecture can be extracted by replacing each *mixed operation* with the most likely one, i.e.  $o^{(i,j)} = \arg \max_{o \in \mathcal{O}} \alpha^{(i,j)}$ .

##### 3.1.2 Optimization

The goal of DARTS is to jointly learn the architecture  $\alpha$  and the weights  $w$  within all the mixed operations. While in RL based methods the performances onto the validation set are used as rewards to train a controller, DARTS is able to optimize the validation loss using gradient descent.

Let  $\mathcal{L}_{train}$  and  $\mathcal{L}_{val}$  be the training and validation loss respectively. Both losses depend both on current architecture  $\alpha$  and on weights  $w$  in the cell operations and nodes. DARTS aims to find an

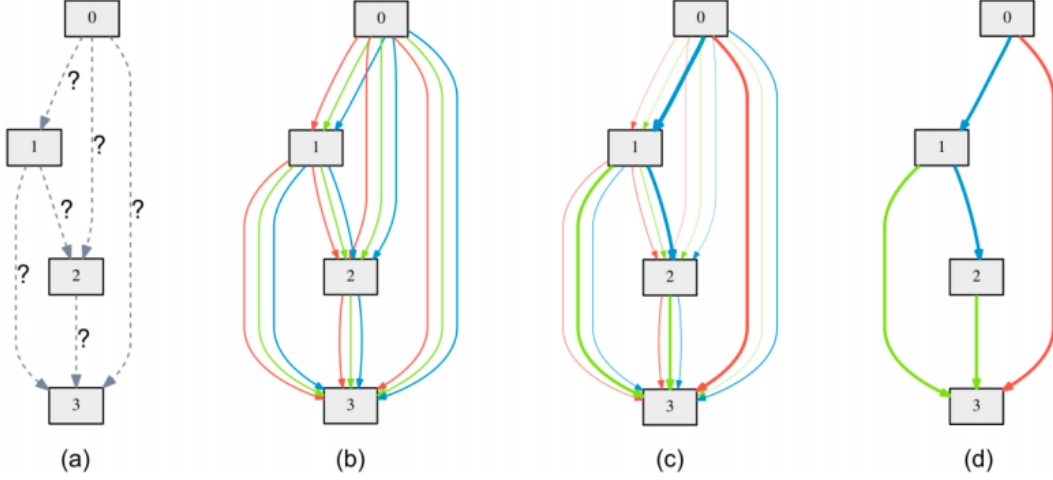


Figure 5: Overview of DARTS, as presented in [15]: Operations on the edges are initially unknown. (b) Continuous relaxation of the search space by placing a mixture of candidate operations on each edge. (c) Joint optimization of the mixing probabilities and the network weights by solving a bi-level optimization problem. (d) Inducing the final architecture from the learned mixing probabilities.

$\alpha^*$  that minimizes the validation loss  $\mathcal{L}_{val}(w^*, \alpha^*)$ , where  $w^* = \arg \min_w \mathcal{L}_{train}(w, \alpha^*)$  are the weights associated to the architecture and are obtained by minimizing the training loss.

It is clear that the optimization of  $\alpha$  depends on the optimization of  $w$ , implying that DARTS optimization task is actually a bi-level optimization problem where  $\alpha$  is the upper-level variable and  $w$  is the lower-level one.

$$\min_{\alpha} \mathcal{L}_{val}(w^*(\alpha), \alpha) \text{ s.t. } w^*(\alpha) = \arg \min_w \mathcal{L}_{train}(w, \alpha) \quad (3)$$

Technically the inner optimization could be solved only by training current weights  $w$  until convergence, but this is only approximated by doing just a single training step, meaning that  $\nabla_{\alpha} \mathcal{L}_{val}(w^*(\alpha), \alpha) \approx \nabla_{\alpha} \mathcal{L}_{val}(w - \xi \nabla_w \mathcal{L}_{train}(w, \alpha), \alpha)$ , with  $\xi$  being the learning rate of the inner optimization. Then upper and lower optimizations are performed in alternated order, as summarized in algorithm 1. The final cell architecture is stacked for a predetermined number of time to form a larger and complete neural network that is then trained and tested.

---

**Algorithm 1:** DARTS – Differentiable Architecture Search

---

Create a mixed operation  $\bar{o}^{(i,j)}$  parametrized by  $\alpha^{(i,j)}$  for each edge  $(i, j)$   
**while** not converged **do**  
    1. Update architecture  $\alpha$  by descending  $\nabla_{\alpha} \mathcal{L}_{val}(w - \xi \nabla_w \mathcal{L}_{train}(w, \alpha), \alpha)$   
    2. Update weights  $w$  by descending  $\nabla_w \mathcal{L}_{train}(w, \alpha)$   
**end**  
Derive the final architecture based on the learned  $\alpha$ .

---

During their experiments, following what Zoph et al. did in 2018 for NASNet [19], two kinds of different cells are searched: *normal cells* and *reduction cells*, with the only difference that operations on reduction cells have stride 2 and are typically placed at 1/3 and 2/3 of the total depth of the network. The architecture encoding  $\alpha$  then becomes a pair of sets  $(\alpha_{normal}, \alpha_{reduce})$ , where  $\alpha_{normal}$  is shared just between normal cells while  $\alpha_{reduce}$  is shared only among reduction cells.

Moreover, DARTS has been able to automatically build both convolutional and recurrent cells, reaching state of the art performances in image recognition and language modeling tasks. This by being orders of magnitude faster than most of currently available NAS systems, with the exception of ENAS, which has comparable performances and computation time.



### 3.2 ENAS

In 2018 H. Phame, B. Zoph et al. proposed a novel neural architecture search approach tailored at fast and inexpensive automatic model design, which goes by the name of ENAS [20]. It actually shares with DARTS most of the aims, like turning the navigation through diverse large search spaces into a feasible operation and avoiding losing weights on the candidate architectures, i.e. *child models*, as soon as they’ve been trained to start over from scratch. It is a controller-based implementation, where the controller, in the form of a recurrent neural network (typically a LSTM), uses the performance of its samples from previous steps as a guiding signal to find more promising architectures.

The main contribution from its introduction is that it pushed the boundaries of NAS systems training without massively dedicated hardware. Indeed, despite being a few % points shy of many human-made networks, it is on-par with other NAS systems performance, while dramatically outperforming them in terms of computational time: depending on the task, ENAS scored a time gain between 3 to 5 orders of magnitude, as it was able to traverse the local regions within a provided macro-architecture space in less than 18 hours on a single GeForce GTX 1080Ti, which go down to a stunning 10 hours for language modeling on the Penny Treebank dataset (Marcus et al., 1994).

#### 3.2.1 Search space

The core idea behind ENAS is the observation that all of the graphs which NAS ends up iterating over can be viewed as sub-graphs of a larger graph (also know in the literature as *super-network*). In other words, we can represent NAS’s search space using a single directed acyclic graph (DAG), whose sub-graphs map generic architectures that could be sampled out of it. In other words, ENAS’s DAG can be seen as the superposition of all its valid child models in the search space, where nodes represent local computations (i.e., convolutional layer, max or average pooling, etc.), each with their own parameters, and edges represent the flow of information (see Figure 6). The parameters at node-level computations’ are used only when that specific computation is activated, which in turn allows ENAS’s children to share weights in the search space.




Figure 6: The above DAG represents the entire search space while the red arrows define a model in the search space, which is decided by a controller. Here, node 1 is the input to the model whereas nodes 3 and 6 are the model’s outputs.

#### 3.2.2 Designing layers

The design of layers changes depending on whether ENAS is trying to produce convolutional or recurrent neural networks, but the baseline task for the controller remains the same. Indeed at each time step (assuming we are at a certain hidden state  $h_i$  within the LSTM) he has two decisions to take: select which nodes should be connected to  $h_i$ -th layer as inputs, and fix the activation function or the local computation, for recurrent and convolutional networks, respectively. Being so similar in nature, and having mentioned both for the sake of completeness, we will just present how ENAS’ controller designs recurrent networks. Let us define  $N = 4$  the number of blocks of decision the LSTM has to go through, in a limited setting, where  $x_t$  is the input signal for a recurrent cell and  $h_{t-1}$  the hidden state from the previous step, as you can see from Figure 7.

- First, the controller samples an activation function that helps it start out from  $h_0$  defining the current state  $h_t = \tanh(x_t \cdot W^{(x)} + h_{t-1}W_1^{(h)})$  where each weight matrix  $W_{(i)}$  has to be learnt and shared across epochs.




Figure 7: An example of a recurrent cell in our search space with 4 computational nodes. Left: The computational DAG that corresponds to the recurrent cell. The red edges represent the flow of information in the graph. Middle: The recurrent cell. Right: The outputs of the controller RNN that result in the cell in the middle and the DAG on the left. Note that nodes 3 and 4 are never sampled by the RNN, so their results are averaged and are treated as the cell’s output.

- For any other node in the graph, the controller samples a previous index (which will function as input) and an activation function. By choosing always a previous index, they are forcing which parameter matrices can be used, because given a pair of indexes  $j, l$  with  $j < l$  there will always be an independent matrix  $W_{l,j}^{(h)}$  since  $l$  won’t ever function as input to  $j$ . Therefore, in ENAS all recurrent cells in a space share the same set of parameters, which is a testament to parameter sharing not being a bad idea, at all.
- For the output, they simply average all the loose ends, i.e., the nodes that didn’t get selected as input to any other node ( $h_3$  and  $h_4$  in the above example).

The search space defined as such includes an exponential number of configurations, which, by assuming 12 nodes ( $N$ ) and 4 plausible activation functions for each of them, returns  $4^N \times N!$  or, approximately,  $10^{15}$  models. Similar numbers can be achieved in the case of convolutional layers, despite them allowing for skip connections through the prediction of previous indexes, via concatenation on depth of inputs, in case one layer has more than one.

### 3.2.3 Optimization

With the above definition of ENAS the optimization unrolls on a bi-level step, where the training of the controller is alternated to training of the sampled architectures. Also, the controller, which samples through softmax classifiers, evolves in an autoregressive fashion as the the decisions it took at the previous step will be fed as its input embedding into the next step; this means that at the first iteration the LSTM will receive an empty input, compliant with the idea that it has absolutely no clue of how a good or bad architecture should look like. There are then two sets of learnable parameters to take into account:  $\theta$ , the controller parameters, and  $w$ , the shared parameters of the child models. The controller will perform an alternating stochastic gradient descent (SGD) by fixing one to descent the other, and viceversa. In particular at first the controller will fix the policy  $\pi(m; \theta)$  and perform SGD on  $w$  to minimize the expected loss  $\mathbb{E}_{m \sim \pi}[\mathcal{L}(m; w)]$ , which is typically chosen as the cross-entropy loss, where  $m$  is a sample drawn from the distribution  $\pi$ . The gradient descent is computed using the Monte-Carlo estimate, approximating the expected loss with the empirical loss on the sequence of drawn models emulating the true distribution:

$$\nabla_w \mathcal{E}_{m \sim \pi}(m; \theta)[\mathcal{L}(m; w)] \approx \frac{1}{M} \sum_{i=1}^M \nabla_w \mathcal{L}(m_i, w) \quad (4)$$

The above result provides an unbiased estimate of the expectation gradient, but still introduces higher variance than standard SGD with  $m$  fixed. By experimenting, they found out an interesting result, that is, the above descent performs pretty well even if  $M = 1$ , i.e. a single sample  $m$  is used to update the gradient, even though they didn’t have a clear answer yet to why that’s the case.  $w$  is then updated through an entire run of the training data on the drawn sample  $m$ . When  $w$  has been updated, the other half of the bi-level optimization takes place: by fixing  $w$  the model tries to update the policy parameters  $\theta$  following the policy gradient in order to maximize the expected reward  $\mathcal{E}_{m \sim \pi(m; \theta)}[\mathcal{R}(m, w)]$  for some reward function  $R$ <sup>4</sup>; they suggest to evaluate the goodness of  $R$

<sup>4</sup>They use REINFORCE algorithm, so they would actually need an intermediate step to compute the cumulative discounted reward

*validation set* rather than on the training set as a proxy to encourage ENAS to look for models in the search space that generalize well. In order to actually assess the best model, the one which scores the higher reward is taken and re-trained from scratch for 150 epochs. Of course, by retraining all the  $M$  models the table of rewards could vary by a good margin, but considering only the one with the initial highest reward is a good greedy approach that gave them a nice balance between accuracy and computational time, as ENAS has actually discovered a novel architecture on the Penny Treebank dataset, with state-of-the-art capabilities.

Extensive experiments have shown super promising results for ENAS, as it scored a 2.89% test error against DenseNet’s 2.56 on CIFAR-10, while comprising less than half of their parameters. Also, the importance of the controller being trained at the same time to the child models weights  $w$  has been assessed (by preventing the controller from training the performance would decrease by up to 5%), as well as the tendency of ENAS to always look for a local minimum in the macro region at its disposal. As we’ve seen, as opposed to DARTS, ENAS can actually deal with both micro- and macro-search with good performance, since a computational DAG, albeit different, is used when local cells need to be predicted instead of the overall architecture, as well; despite that, how you would choose the macro-architecture to feed it with, in the case of a micro-search setting for ENAS, doesn’t have a precise answer yet. All in all, ENAS has left its mark on NAS research since it has paved the way to a lot of the current exploration in finding efficient ways to deal with macro-search assessment.

## 4 Proposed implementation

We provide as a part of this survey a naive implementation of a controller-based NAS system<sup>5</sup>, whose main objective is sampling multi-layer perceptrons (MLPs) capable of solving a given task. This can be seen as a toy problem, since much more complex models should be sampled to tackle real-world tasks, for example CNNs, and those would need a much deeper reasoning on how to structure the search space, how to describe every single convolutional layer, and so on. The choice of searching for MLPs also convinced us to use a macro-search approach, as cell searching would have been a useless complication, given the scope of the implementation.

This NAS system is based on a LSTM model whose goal is to sample a probability distribution for every possible layer accounted by the search space, and then extract the next predicted layer through a random choice based on this distribution. At every epoch, the controller calls the MLP generator for a number of times equal to the desired number of architectures to sample per-epoch. The layers of these architectures are sampled in terms of unique indexes compliant to a vocabulary, that maps to a specific pair in the form of *(n. of neurons, activation function)*.

When all of the layers for a single architecture have been sampled, the whole architecture is decoded back to actual Keras layers, which are concatenated in order to obtain a trainable model. After the sampling phase, each architecture is trained on its dataset of choice for approximately 10 epochs with an early stopping of 3. This training procedure can actually take quite some time depending on the complexity of the dataset, so this must be one of the deciding factors of both the number of sampling epochs the controller performs and the number of architectures being sampled per-epoch.

### 4.1 Design choices

Many of the design choices we made were based primarily on the original NAS paper [21] and its more recent evolution, ENAS [20]. Alongside the classical RL based NAS mechanics, we also implemented a one-shot architecture search feature, meaning that identical pairs of layers (also addressed as *bigrams*) share weights connecting them. In order to do so we use a dictionary that stores bigrams of the layers sampled so far and their corresponding weights, that will be transferred to newly sampled architectures whenever the opportunity arises.

The code should be self-explanatory and structured in modular Python files, but we will still comment on a few key notes worth of mention. Regarding the controller, more LSTM layers could have been implied to sample the probability distributions, but we didn’t test it out for time constraints. A few questions which may actually sound more appealing are what the true role of the NAS’ controller is and how the LSTM itself would actually be trained. Its main role is to assure that the sequences

---

<sup>5</sup>Source code can be found on the repository at <https://github.com/ResonantFilter/AMLProject-NAS>

being sampled by the MLP generator are valid and that no repeating ones are being sampled to avoid wasting important computational resources on regions of the search space which have already been explored. But how well the controller traverses the search space depends on how well it is trained to follow the more optimal directions. After the first few sequences have been generated, trained and evaluated they will form a dataset the controller is fed with. In essence, after every controller epoch, a new dataset that collects the architectures sampled since the first epoch is created for the controller to learn from, and this should make it learn how to discriminate within its hidden states between architectures that perform well from those that do not.

The actual training of the controller must be dependant from the validation accuracies obtained by its generated architectures, thus its loss function implements the aforementioned REINFORCE algorithm. In our case, the controller functions as a *model-free agent*, that is an agent that doesn't try to understand the environment, but rather just focuses on the policy  $\pi$ .

## 4.2 Accuracy predictor

Our implementation opens up for the possibility to introduce an accuracy predictor, a network parallel to the LSTM itself, which turns the model into an adversarial model by accounting for an optimization that doesn't focus solely on the above REINFORCE loss function, but also on how well the accuracy predictor is becoming in predicting the goodness of the sampled architectures on their task. Indeed the predictor is implemented by using a single dense layer which will share weights with the LSTM layer of the sequence layer, fed with the true validation accuracies reported by the previously sampled architectures at time  $t$ , thus letting it construct an internal representation of the architectures that allows it to understand the properties that characterize a good architecture as opposed to a bad one, without the need to train them for those 10 epochs, as a proxy to their validation accuracies. The controller, on the other hand, will try to navigate the search space in a way that also allows it to generate architectures not easily predictable by the predictor.

Despite the adversion of the predictor may lead to architectures with lower validation accuracies on some tasks, its usage is still desired to help sampling architectures that generalize better.

# 5 Results

## 5.1 Data and configurations

We conducted our experiments on three datasets:

1. MNIST Digits [22]
2. Wine-Quality
3. 10 Speakers, a dataset that has been used as part of a project for "Speech Processing and Recognition" course, containing audio features from 10 different speakers.

In particular we did extensive testing using "MNIST Digits" dataset, while "Wine-Quality" and "10 speakers" were included as additional playgrounds. Then we defined a search space for MLPNAS that can be summarized as follows:

- Possible layer widths: [8, 16, 32, 64, 128, 256, 512]
- Possible activation functions: [sigmoid, tanh, ReLU, ELU]
- Maximum architecture length: 10 layers

We also included the chance of sampling dropouts layer, for which dropout rate has been set constant to 0.2, for the sake of simplicity. The final layer is decided depending on the given task: we use a single neuron sigmoid in case of binary classification problems, while multi-classification uses a final softmax layer whose width is automatically set accordingly w.r.t labels inside the dataset.

Many hyperparameters for both the controller and generated architectures could have been tuned, but we could not do that in depth due to time limitations. We train the LSTM controller for 10 epochs, and for each epoch 10 architectures are generated. Each of the generated architectures are trained for 10 epochs as well. Both controller and child networks trainings are optimized using

Adam optimizer [23] with learning rates fixed at 0.01 and 0.001 for the controller and generated architectures respectively, keeping Keras' default value for  $\beta$ . The controller uses a learning rate decay of 0.1, while child networks use none.

Training time with LSTM's hidden size of 512 and looking for 20-layer-long architectures at maximum, took around 1 hour on a GeForce GTX 1660 Ti GPU on MNIST Digits dataset. The default LSTM's size has been fixed to 128, though, as the data we tested out was quite simple in nature (only numerical or categorical features), thus our MLPNAS model didn't need a bigger hidden space to learn a powerful enough input-output function  $\Phi$ .

We also made an experiment in which we increased to 15 the number of epochs for the Controller as well as the number of architectures sampled per epoch, that has been brought to 30. With this configuration we encountered many promising architectures, even if it must be pointed out that we are dealing with a toy problem on a very simple dataset, and complex architectures may not be worthy. As we show in Table 4, even if the first architecture is the best we found so far with a 96% validation accuracy, it is too complex for the task at hand. This is clear when we look at the second architecture in Table 4, which reaches a 95% validation accuracy with just 2 layers.

## 5.2 Obtained MLP architectures

Tables in this section display architectures found by MLPNAS. In particular we show the worst and the best architectures among top 5 layouts discovered by MLPNAS for each task. *Please notice that validation accuracy refers to validation accuracy obtained after 10 epochs of training during architecture generation, while test accuracy is obtained on a test set after 50 additional epochs of fine-tuning since we did a full train only on the most promising architecture.*

For each result, we also provide the test accuracy obtained from fitting a Scikit-learn's MLP Classifier [24] onto the same dataset.

Table 1: **MNIST Digits** – Without one-shot learning and predictor on 128 hidden size LSTM

Architecture	Validation Accuracy	Test Accuracy
Scikit-learn MLP test accuracy	–	0.93
(256, 'sigmoid'), (8, 'elu'), (256, 'tanh'), (256, 'relu'), (10, 'softmax')	0.68	–
<b>(32, 'elu'), (512, 'sigmoid'), (32, 'elu'), (512, 'sigmoid'), (10, 'softmax')</b>	<b>0.72</b>	<b>0.72</b>

Table 2: **MNIST Digits** – Using predictor only on 128 hidden size LSTM

Architecture	Validation Accuracy	Test Accuracy
Scikit-learn MLP test accuracy	–	0.91
(128, 'sigmoid'), (64, 'sigmoid'), (16, 'relu'), (512, 'elu'), (10, 'softmax')	0.66	–
<b>(32, 'elu'), (32, 'relu'), (64, 'relu'), (64, 'elu'), (10, 'softmax')</b>	<b>0.74</b>	<b>0.86</b>

Table 3: **MNIST Digits** – Using one-shot learning and predictor on 128 hidden size LSTM, controller batch size set to how many architectures are sampled per epoch.

Architecture	Validation Accuracy	Test Accuracy
Scikit-learn MLP test accuracy	–	0.92
(64, 'sigmoid'), (512, 'relu'), (10, 'softmax')	0.70	–
<b>(512, 'sigmoid'), (10, 'softmax')</b>	<b>0.77</b>	<b>0.84</b>

Table 4: **MNIST Digits** – Using one-shot learning and predictor on 128 hidden size LSTM, with extended controller batch size and amount of sampled architectures per epoch.

Architecture	Validation Accuracy	Test Accuracy
Scikit-learn MLP test accuracy	–	0.92
(128, 'elu'), (64, 'relu'), (64, 'relu'), (32, 'relu'), (512, 'sigmoid'), (64, 'tanh'), (64, 'relu'), (32, 'relu'), (256, 'relu'), (10, 'softmax')	<b>0.96</b>	–
(256, 'elu'), (10, 'softmax')	0.95	–

Table 5: **10 Speakers** – Using one-shot learning and predictor on 128 hidden size LSTM, controller batch size set to how many architectures are sampled per epoch.

Architecture	Validation Accuracy	Test Accuracy
Scikit-learn MLP test accuracy	–	0.21
(64, elu), (32, relu), (16, tanh), (512, relu), (32, elu), (16, sigmoid), (64, elu), (16, elu), (512, relu), (10, softmax)	0.42	–
(128, elu), (64, elu), (32, relu), (128, sigmoid), (32, relu), (8, elu), (32, relu), (16, elu), (32, tanh), (10, softmax)	<b>0.43</b>	<b>0.47</b>

Table 6: **Wine Quality** – Using one-shot learning and predictor on 512 hidden size LSTM

Architecture	Validation Accuracy	Test Accuracy
Scikit-learn MLP test accuracy	–	0.69
(512, sigmoid), (512, relu), dropout, (512, relu), (256, elu), (256, elu), (256, tanh), (512, elu), (256, elu), (3, softmax)	0.64	–
(16, sigmoid), (512, sigmoid), (3, softmax)	<b>0.66</b>	<b>0.64</b>

## 6 Conclusions

Although NAS methods are constantly improving, the quality of empirical evaluation in this field is still lagging behind when compared to other areas in machine learning, AI and optimization. Hence, more efficient performance evaluation strategies are to be researched, in order to give NAS systems a true shot at competing against human-made networks. While this branch of AutoML has clear potential in terms of lifting deep learning experts from the task of manually designing architectures, we think that – more importantly – studying in depth the properties of good neural networks gives us the chance of truly understand them, in a way that is not possible to this day.

## References

- [1] Dr. Debadeepta Dey. Advanced machine learning day 3: Neural architecture search, October.
- [2] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. 2015.
- [3] Frank Hutter and Feurer M. Hyperparameter optimization, in automated machine learning. *Automated Machine Learning*, 2019.
- [4] Frank Hutter, Lars Kotthoff, and Joaquin Vanschoren. Automated machine learning: Methods, systems, challenges. *Automated Machine Learning*, 2019.
- [5] Barret Zoph and Quoc V. Le. Neural architecture search with reinforcement learning. *CoRR*, abs/1611.01578, 2016.
- [6] Barret Zoph et al. <https://ai.googleblog.com/2017/11/automl-for-large-scale-image.html>.
- [7] <https://analyticsindiamag.com/why-tesla-invented-a-new-neural-network/>.
- [8] [https://nni.readthedocs.io/en/latest/NAS/one\\_shot\\_nas.html](https://nni.readthedocs.io/en/latest/NAS/one_shot_nas.html).
- [9] Thomas Elsken, Jan Hendrik Metzen, and Frank Hutter. Neural architecture search: A survey. *ArXiv*, abs/1808.05377, 2019.
- [10] Hanzhang Hu. Macro neural architecture search revisited. 2019.
- [11] Karim Ahmed and Lorenzo Torresani. Connectivity learning in multi-branch networks. *ArXiv*, abs/1709.09582, 2017.
- [12] Shreyas Saxena and Jakob Verbeek. Convolutional neural fabrics. *ArXiv*, abs/1606.02492, 2016.
- [13] Richard Shin, Charles Packer, and Dawn Xiaodong Song. Differentiable neural network architecture search. In *ICLR*, 2018.
- [14] Tom Veniat and Ludovic Denoyer. Learning time-efficient deep architectures with budgeted super networks. *ArXiv*, abs/1706.00046, 2017.
- [15] Hanxiao Liu, Karen Simonyan, and Yiming Yang. Darts: Differentiable architecture search. *ArXiv*, abs/1806.09055, 2019.
- [16] <https://ai.googleblog.com/2019/04/morphnet-towards-faster-and-smaller.html>.
- [17] George Kyriakides and Konstantinos G. Margaritis. The effect of reduced training in neural architecture search. *Neural Computing and Applications*, 04 2020.
- [18] Jeovane Honorio Alves and Lucas Ferrari de Oliveira. Optimizing neural architecture search using limited gpu time in a dynamic search space: A gene expression programming approach, 2020.
- [19] Barret Zoph, V. Vasudevan, Jonathon Shlens, and Quoc V. Le. Learning transferable architectures for scalable image recognition. *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 8697–8710, 2018.

- [20] Hieu Pham, Melody Y. Guan, Barret Zoph, Quoc V. Le, and Jeff Dean. Efficient neural architecture search via parameter sharing. *CoRR*, abs/1802.03268, 2018.
- [21] Barret Zoph and Quoc V. Le. Neural architecture search with reinforcement learning. *ArXiv*, abs/1611.01578, 2017.
- [22] Yann LeCun and Corinna Cortes. MNIST handwritten digit database. 2010.
- [23] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization, 2014. cite arxiv:1412.6980Comment: Published as a conference paper at the 3rd International Conference for Learning Representations, San Diego, 2015.
- [24] Scikit-learn mlp classifier. [https://scikit-learn.org/stable/modules/generated/sklearn.neural\\_network.MLPClassifier.html](https://scikit-learn.org/stable/modules/generated/sklearn.neural_network.MLPClassifier.html).

## 7 Appendix

### 7.1 REINFORCE algorithm

A loss often implied to maximize the policy gradient is given by the REINFORCE algorithm, which is a typical algorithm in RL that represents a Monte-Carlo variant of a stochastic policy gradient. Its objective is then to learn a policy that maximizes the cumulative future reward score  $G$  by means of a combination of the rewards produced by a reward function  $R$ , computed on all the  $(a_i, s_i)$  pairs, where the policy  $\pi$  is defined as a probability distribution of actions in which a higher probability value corresponds to a higher expected reward for an action from a given observed state.

The REINFORCE algorithm tries to maximize an objective  $J$  defined as the product of the cumulative future discounted reward  $G(t)$  w.r.t. a baseline (0.5, for example, that basically will make it discriminate about half of the actions as good and half as bad) with log-probabilities of actions w.r.t. the policy as  $\ln \pi(A|S, \theta)$ , where  $A = [a_1, \dots, a_T]$ . Formally, it can be written as:

$$J(\theta) = \alpha \gamma^t G [\nabla_{\theta} \ln \pi(a_t | s_t, \theta)]$$

---

#### Algorithm 2: REINFORCE

---

```

N number of epochs;
T number of steps;
;
θ random;
;
for  $n = 1, \dots, N$  do
    Generate an episode,  $(s_1, a_1, r_1), \dots, (s_T, a_T, r_T)$ , following the current policy  $\pi(\cdot | \cdot, \theta)$ ;
    for  $t = 1, \dots, T$  do
        Compute the discounted cumulative reward  $G(t)$  6;
         $\theta \leftarrow \theta + \alpha \gamma^t G [\nabla_{\theta} \ln \pi(a_t | s_t, \theta)]$ ;
    end
end

```

---

where  $\gamma$  is the discount factor and  $\pi(a_t | s_t, \theta)$  the probability of the occurrence of  $(s_t, a_t)$  given the current trajectory  $\tau$  followed by the agent, regulated by the learning rate,  $\alpha$ . Indeed, the learning rate is a key factor in traversing appropriate regions of the search space, because a NAS model could very easily get stuck in regions where no good architectures can be found for the problem at hand.

Since in REINFORCE algorithm the sample gradient expectation is equal to the actual gradient, expressed in terms of log-probabilities on  $\pi$ , it reflects a good theoretical convergence, albeit being a Monte-Carlo-based method it may suffer from high variance.

---

<sup>6</sup>Please, refer to *mlpnas.py* in the code for details on  $G(t)$  computation.